

CSEE 3827: Fundamentals of Computer Systems, Spring 2022

Lecture I

Prof. Dan Rubenstein (danr@cs.columbia.edu)

Quick Announcements

- Please fill in Office Hour Availability at:
- <http://uribe.cs.columbia.edu/sched/table.php> (due Tues Jan 25)
- Waitlisted students:
 - Section 1 & 2: please enroll in Section 3 for now (1 & 2 overloaded on waitlist)
 - Section 3: larger room forthcoming... will assign as soon as possible
 - EdStem bug for “Observers” - reported the bug, they’ll try to fix... in the meantime I’m adding folks manually

A note on lecture slides

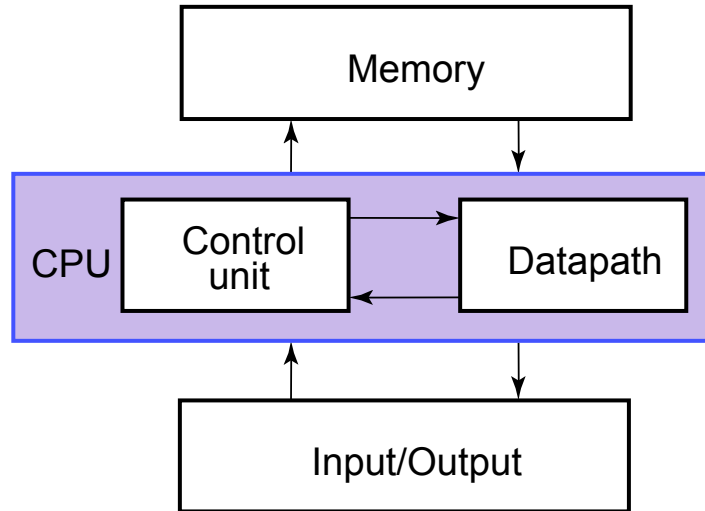
- There are 13 lecture slides over the course of the term
- They are grouped by topic
- Obviously, we won't always complete a set of lecture slides in a single lecture. Some can take 2,3, or 4 class lectures

Agenda (M&K&M Ch 1, 3.11, 9.7)

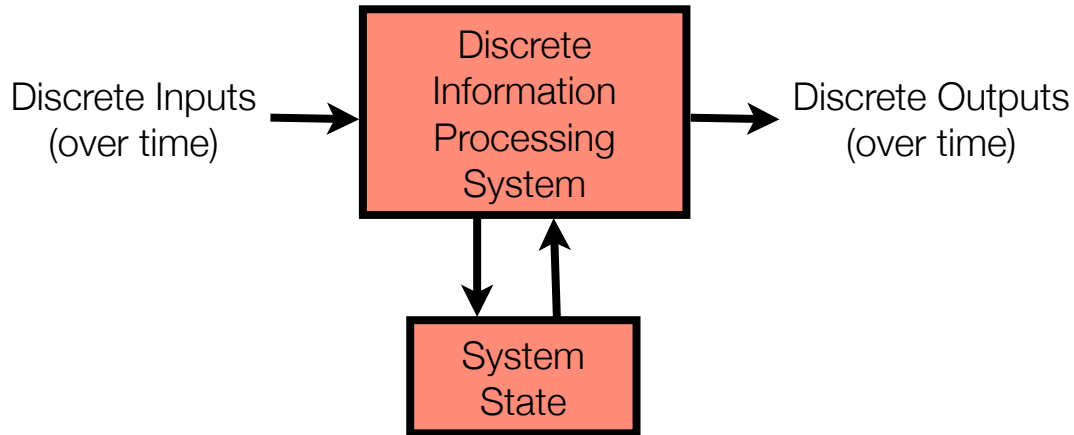
- Computer from a (very high level) Digital Perspective
- Digital/Binary
 - vs Decimal, Hexadecimal
- Terminology:
 - Bit / Byte / **Word & Wordsize**
 - **Highest Order (most significant) Bit, Lowest Order (least significant) bit**
- Negative Number Formats:
 - Signed Magnitude
 - 1's Complement
 - **2's Complement**
- Floating Point via Binary (P&H 3.5 - skip FP in MIPS subsection)
 - Addition
 - Multiplication

Course Overview: Building a computer (digital perspective)

- **CPU**: the “brain” of a computer
 - **Control unit** does calculations on data in **datapath**
- **Memory**: stores data (for later use)
- **Input/Output**: interface to outside (disk, network, monitor, keyboard, mouse, etc.)

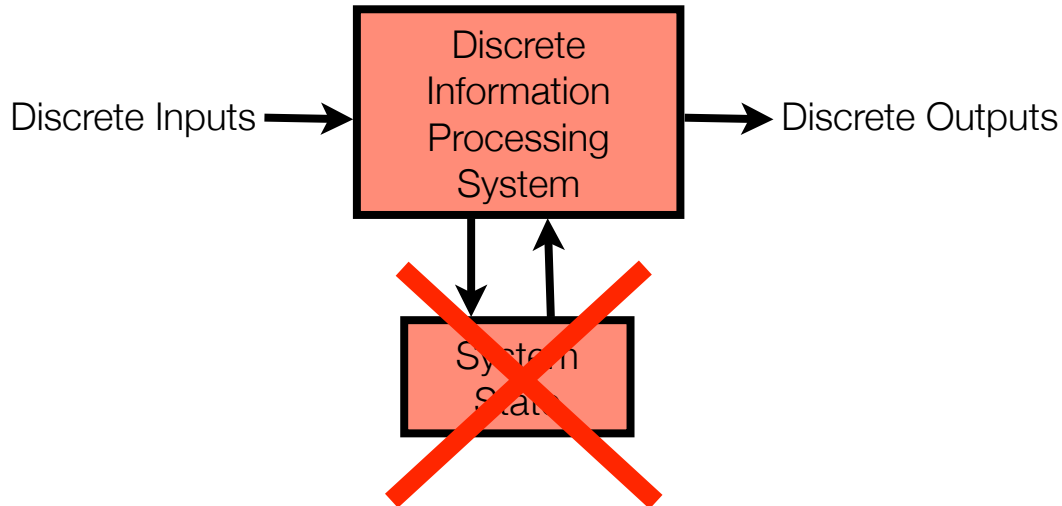


More simplistic view



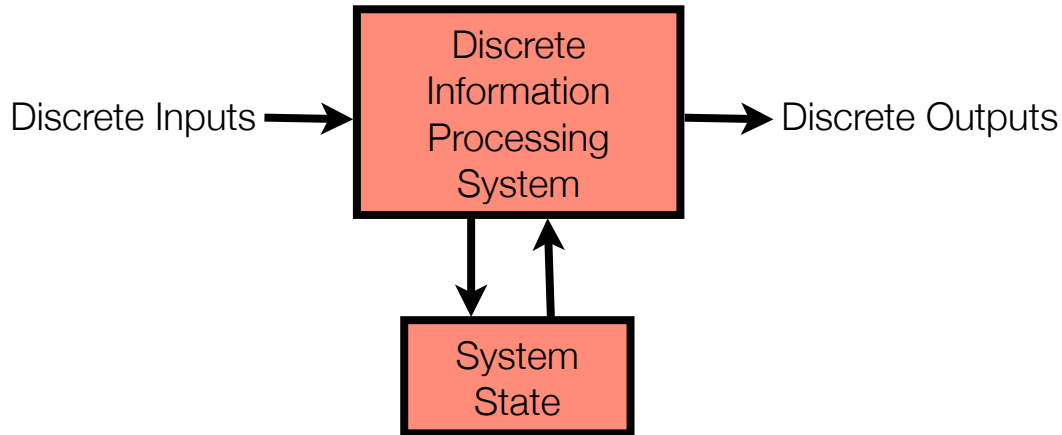
- Course: Starts with this more simple view (on next slide)

Course Overview: 1st quarter



- 1st quarter of course: really simple view: “computer” doesn’t maintain state
- Input → Compute → Output (just a math function)
- Feed same input, get same output

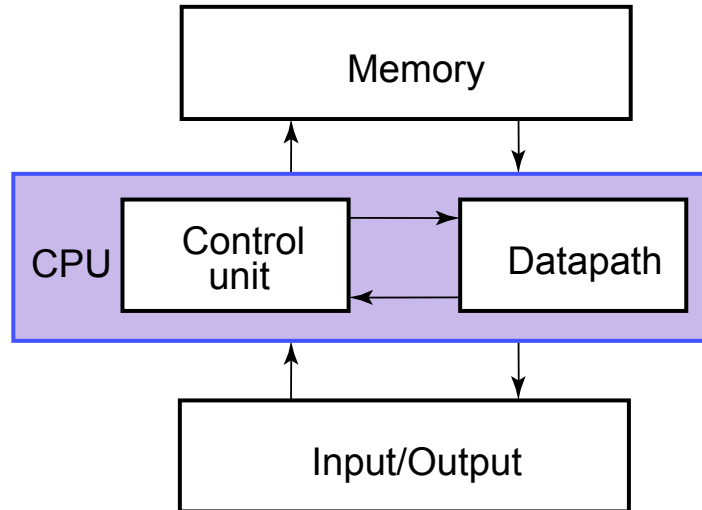
Course Overview: 2nd quarter



- 2nd quarter: “computer” has memory (system state)
 - Can use input & what it has stored in memory to determine output

Course Overview: 2nd Half

- Computer processes **programs** (stored in memory)
 - program made up of sequences of **instructions**
- Programs modify data also stored in memory



More on this later in term...

Addition in different bases

Number systems review: Base 10 (Decimal)

- How humans (usually) work with numbers
- 10 digits = $\{0,1,2,3,4,5,6,7,8,9\}$
- example: 4537.8 base 10 a.k.a. $(4537.8)_{10}$

$$\begin{array}{ccccccccc} 4 & & 5 & & 3 & & 7 & & . & 8 \\ \times 10^3 & & \times 10^2 & & \times 10^1 & & \times 10^0 & & \times 10^{-1} & \\ \hline 4000 & + & 500 & + & 30 & + & 7 & + & .8 & = 4537.8 \end{array}$$

Number systems review: Base 10 (Decimal)

- How humans (usually) work with numbers
- 10 digits = $\{0,1,2,3,4,5,6,7,8,9\}$
- example: 4537.8 base 10 a.k.a. $(4537.8)_{10}$

$$\begin{array}{ccccccccc} 4 & & 5 & & 3 & & 7 & & . & 8 \\ \times 10^3 & & \times 10^2 & & \times 10^1 & & \times 10^0 & & \times 10^{-1} \\ \hline 4000 & + & 500 & + & 30 & + & 7 & + & .8 & = & 4537.8 \end{array}$$

Shifting a column to the left multiplies a digit's value by 10

Number systems: Base 2 (Binary)

- How computers “think” about numbers
- 2 digits = {0,1}
- example: $(1011.1)_2$

$$\begin{array}{rccccccccc} & 1 & & 0 & & 1 & & 1 & . & 1 \\ & \times 2^3 & & \times 2^2 & & \times 2^1 & & \times 2^0 & & \times 2^{-1} \\ \hline & 8 & + & 0 & + & 2 & + & 1 & + & .5 & = (11.5)_{10} \end{array}$$

Number systems: Base 2 (Binary)

- How computers “think” about numbers
- 2 digits = {0,1}: we refer to binary digits as **bits**
- example: $(1011.1)_2$

$$\begin{array}{rccccccccc} & 1 & & 0 & & 1 & & 1 & & . & 1 \\ & \times 2^3 & & \times 2^2 & & \times 2^1 & & \times 2^0 & & \times 2^{-1} \\ \hline 8 & + & 0 & + & 2 & + & 1 & + & .5 & = (11.5)_{10} \end{array}$$

Shifting a column to the left multiplies a digit's value by 2

Number systems: Base 16 (Hexadecimal)

- How (nerdy) humans interact with computers
- 16 digits = {0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F}

Base 16 (Hexadecimal) Value	Base 2 (binary) value	Base 10 value
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
A	1010	10
B	1011	11
C	1100	12
D	1101	13
E	1110	14
F	1111	15

Number systems: Base 16 (Hexadecimal)

- How (nerdy) humans interact with computers

- 16 digits = {0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F}

- example: (26BA) [alternate notation for hex: 0x26BA]

$$\begin{array}{cccc} 2 & 6 & B & A \\ \times 16^3 & \times 16^2 & \times 16^1 & \times 16^0 \\ \hline 8192 & + & 1536 & + & 176 & + & 10 & = & (9914)_{10} \end{array}$$

Number systems: Base 16 (Hexadecimal)

- How (nerdy) humans interact with computers

- 16 digits = {0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F}

- example: (26BA) [alternate notation for hex: 0x26BA]

$$\begin{array}{cccc} 2 & 6 & B & A \\ \times 16^3 & \times 16^2 & \times 16^1 & \times 16^0 \\ \hline \end{array}$$

$$8192 + 1536 + 176 + 10 = (9914)_{10}$$

- Shifting a column to the left multiplies a digit's value by $(16)_{10}$
- Why Important: More concise than binary, but related (a power of 2)

Number ranges

- Cannot map infinite numbers in bounded memory: restrict to some finite range: must choose the representation for a computer
- How many numbers can I represent with ...

... 5 digits in decimal? 10^5 possible values $(100,000)_{10}$

... 8 binary digits? 2^8 possible values $(256)_{10}$

... 4 hexadecimal digits? 16^4 possible values $(65,536)_{10}$

Why Base 16 is useful

- Computers work in base 2 (binary)
- Writing base 2 numbers is confusing / laborious
- Easy to convert between base 2 and base 16, here's why:
 - To multiply a base-2 4-bit quantity by 16, shift over 4 columns
 - e.g., 1101 0000 is 16x larger than 1101
 - To multiply a base-16 1-digit quantity by 16, shift over 1 column
 - e.g., D0 is 16x larger than D

Converting between bases 2 and 16

- 26BA
 - = 2000 + 600 + B0 + A
 - = 2 (shifted left 3) + 6 (shifted left 2) + B (shifted left 1) + A (hex digits)
 - = 0010 (shifted left 12) + 0110 (shifted left 8) + 1011 (shifted left 4) + 1010 (binary digits)
 - = 0010 0000 0000 0000 + 0110 0000 0000 + 1011 0000 + 1010
 - = 0010 0110 1011 1010
- 2 6 B A
-
- Diagram illustrating the conversion of the hexadecimal number 26BA to binary:
- HEX
 - Binary

Using Base 16

- Useful for referring to long sequences of binary, e.g., 32-bit quantity:
- 0110 1010 1111 1010 0000 1100 0011 1111
- 6 A F A 0 C 3 F
- i.e., the 32-bit quantity can be written 6AFA0C3F, which is
- 01101010111110100000110000111111

Defs & Terminology

Computer from Digital Perspective


- (Digital) Information: just sequences of binary (0's and 1's)
 - True = 1, False = 0
 - Numbers: converted into binary form when “viewed” by computer
 - e.g., $19 = 10011$ ($16(1) + 8(0) + 4(0) + 2(1) + 1(1)$) in binary
 - Characters: assigned a specific numerical value (ASCII standard)
 - e.g., 'A' = 65 = 1000001, 'a' = 97 = 1100001
 - Text is a sequence of characters:
 - “Hi there” = 72, 105, 32, 116, 104, 101, 114, 101
 - = 1001000, 1101001, ...

Terminology: Bit, Byte, Word

- **Bit**: a single binary digit (a '0' or a '1')
- **Byte**: a grouping of 8 bits, e.g., 10110010. Q: how many distinct bytes exist?
- **Word**: a grouping of bits that is **computer-architecture dependent**
 - The number of bits that the computer architecture can process at once
 - e.g., 64-bit word architectures expect data to be passed in (and returns outputs back out) in 64-bit groupings
 - The number of bits used by the architecture is called its **word size**
 - **OBSERVATION**: computers have bounds on how much input they can handle at once
 - word size limits on the sizes of numbers they can deal with (in a single computation cycle)

Terminology 2: Highest / Lowest Order / significant bits


- Bit at the left is **highest order (a.k.a. most significant)** bit
- Bit at the right is **lowest order (a.k.a. least significant)** bit

- e.g.,  1 0 1 0 1 1 1
- higher order bits represent “bigger” values when read as numbers
- e.g., above is $64 + 16 + 4 + 2 + 1 = 87$ if read as unsigned binary
- Common reference notation for k-bit value: $b_{k-1}b_{k-2}b_{k-3}...b_1b_0$
 - or alternatively: $b_kb_{k-1}b_{k-2}...b_2b_1$

Modular Arithmetic

Modular Arithmetic

- Def: $X \bmod Y = \text{remainder}(X/Y)$ (i.e., a value between 0 and $Y-1$)
- e.g.,
 - $22 \bmod 5 = 2$ ($22 / 5 = 4$ with remainder 2, i.e., $22 = 5 * 4 + 2$)
 - $100 \bmod 9 = 1$,
 - $-10 \bmod 3 = -1 \bmod 3 = 2 \bmod 3$ ($-4 * 3 + 2$)

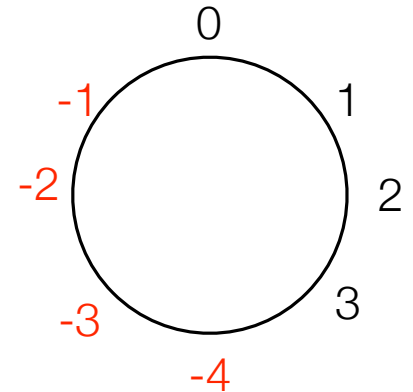


“Pinwheel” representation: move clockwise on the pinwheel for + numbers, counter-clockwise for - numbers

This pinwheel is for mod 8

Modular Arithmetic with negative representations

- For $X \bmod Y$, we usually define the remainder as a value between 0 and $Y-1$
- Alternative: define remainder between $-Y/2$ and $Y/2-1$
- e.g., mod 8 would use remainder values -4 through 3:



“Pinwheel” representation: move clockwise on the pinwheel for + numbers, counter-clockwise for - numbers

- $4 \bmod 8 = -4$ ($4 = 8 \cdot 1 + (-4)$)
- $7 \bmod 8 = -1$ ($7 = 8 \cdot 1 + (-1)$)
- $-5 \bmod 8 = 3$ ($-5 = 8 \cdot -1 + 3$)
- $30 \bmod 8 = -2$ ($30 = 8 \cdot 4 + (-2)$)
- $30 \bmod 10 = 0$ ($30 = 10 \cdot 3 + 0$)

This pinwheel is for mod 8 (range -4 to 3)

Important for 2's complement (discussed later)

Integer # Formats

(with word size restriction)

Suppose word size = k (e.g., $k = 3$)

- How many different bit-sequences are there with word size k ?

- 2^k (e.g., for $k=3$, $2^3 = 8$)

000
001
010
011
100
101
110
111

Suppose word size = k (e.g., k = 3)

- How many different bit-sequences are there with word size k?

- 2^k (e.g., for $k=3$, $2^3 = 8$)

000
001
010
011
100
101
110
111

- Can assign each number value to a distinct bit-sequence: can have up to 2^k different number values

A weird but feasible assignment:

word	val
000	15
001	-7
010	42
011	42
100	7
101	9
110	e^7
111	0.004

Unsigned binary numbers (with wordsize restriction)

- Binary numbers represent only non-negative (positive or 0) values
- e.g., wordsize = 4 (computer “thinks” using 4 bits at a time)
 - 0000 = 0
 - 0011 = 3
 - 1011 = 11
 - 1111 = 15
 - Can’t represent 16 or larger as unsigned binary with wordsize of 4!!
 - Can’t represent negative #'s as unsigned binary

BAA

Binary Addition Algorithm (of unsigned numbers)

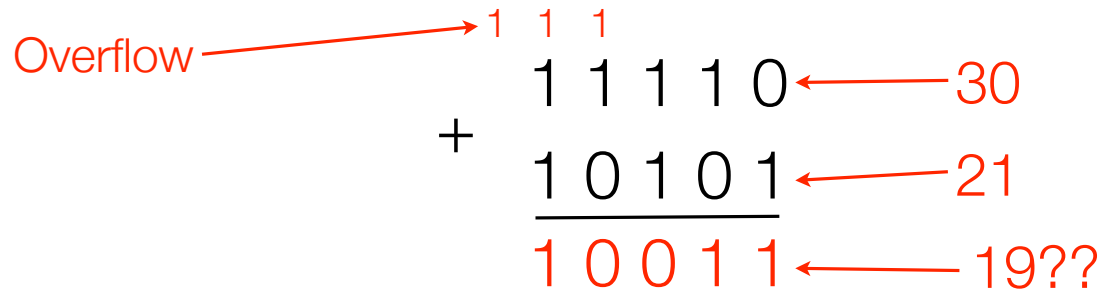
- Like regular (base 10) addition algorithm, except:
 - $1+1 = 0$ with a carry of 1, $1+1+1 = 1$ with carry of 1
 - e.g., wordsize = 5 (all numerical representations restricted to 5 bits)
 - add 11110 and 10101 (30 + 21)

$$\begin{array}{r} 11110 \\ + 10101 \\ \hline \end{array}$$

BAA

Binary Addition Algorithm (of unsigned numbers)

- Like regular (base 10) addition, except:
 - $1+1 = 0$ with a carry of 1, $1+1+1 = 1$ with carry of 1
 - e.g., wordsize = 5, add 11110 and 10101 (30 + 21)



A binary addition diagram for 5-bit unsigned numbers. The first number is 11110 (30) and the second is 10101 (21). The sum is 10011 (19). The result 10011 is shown in red. A red arrow points from the word 'Overflow' to the carry-out '1' above the first bit of the result. Red arrows also point from the decimal values 30, 21, and 19?? to their respective binary representations.

$$\begin{array}{r} \text{Overflow} \rightarrow 1 \ 1 \ 1 \\ 1 \ 1 \ 1 \ 1 \ 0 \leftarrow 30 \\ + \quad 1 \ 0 \ 1 \ 0 \ 1 \leftarrow 21 \\ \hline 1 \ 0 \ 0 \ 1 \ 1 \leftarrow 19?? \end{array}$$

- **Overflow:** when result cannot fit within the wordsize constraint
- e.g., above, the “correct” answer 110011 (=51) requires 6 bits: cannot be represented with only 5 bits in unsigned representation

Modular Arithmetic and Overflow

$$\begin{array}{rcl} & 1 & 1 & 1 & 1 & 0 & 30 \\ + & 1 & 0 & 1 & 0 & 1 & 21 \\ \hline & 1 & 0 & 0 & 1 & 1 & 19?? \end{array}$$

- Suppose we hadn't restricted wordsize, and allowed the solution to use an extra bit
- Solution would have been $110011 = 51$
- The word size restriction prevented us from including the bit = 32
- The result (19) is off by 32, but is correct $\text{mod } 32 = \text{mod } 2^{\text{wordsize}}$
- i.e., $19 = 51 \text{ mod } 2^5$
- In other words, arithmetic within fixed wordsize computes the remainder

Negative Numbers

- Given: computer has a fixed wordsize (e.g., 4)
- Q: How to represent both positive and negative #'s when constrained by this fixed wordsize?
 - Note: we need an alternate mapping (from what is used for unsigned binary) of bit sequences to values so that some bit sequences represent negative numbers
- A: There are several ways to do this, some are easier for humans, some for computers...

Negative

Representations:

- Signed magnitude
- 1's Complement
- 2's Complement

Negative Numbers: Signed Magnitude representation

- highest order bit (b_{k-1}) indicates **sign**: 0 = positive, 1 = negative
- remaining bits indicate **magnitude**
 - e.g., 0011 = 3
 - e.g., 1011 = -3
 - e.g., 1000 = 0000 = 0
- positive #'s have same form in both signed magnitude and unsigned
- Easy for humans to interpret, but not easiest form for computers to do addition/subtraction operations (as we will soon see...)

Negative Numbers: 1's Complement representation

- Again, highest order bit (b_{k-1}) indicates **sign**: 0 = positive, 1 = negative
- Non-negative #'s have same representation as unsigned (and signed-mag)
- To negate a #, **flip all bits** (not just highest-order as in signed-mag)
- Note flipping all bits will also flip the highest order (sign) bit
- e.g., wordsize = 4
 - 0010 = 2
 - 1101 = -2

Negative Numbers: 1's Complement representation

- Non-negative #'s have same representation as unsigned (and signed-mag)
- To negate a #, flip all bits (not just highest-order as in signed-mag)
- e.g., wordsize = 4
 - $0010 = 2$
 - $1101 = -2$
- Q: suppose wordsize is 8, what is the value of 11101011 when it represents a # in 1's Complement representation?
- A: Let $X = 11101011$ (it's a negative #)
 - Negate X by flipping all bits: $-X = 00010100$
 - $-X = 20$, so $X = -20$

Negative Numbers: 1's Complement representation

- Non-negative #'s have same representation as unsigned (and signed-mag)
- To negate a #, **flip all bits** (not just highest-order as in signed-mag)
- e.g., wordsize = 4
 - $0010 = 2$
 - $1101 = -2$
- Q: suppose wordsize is 8, what is the value of 11101011 when it represents a # in 1's Complement representation?
- A: Let $X = 11101011$ (it's a negative #)
 - Negate X by flipping all bits: $-X = 00010100$
 - $-X = 20$, so $X = -20$
- NOTE: 2 ways to represent 0 in 1's Complement: all 0's and all 1's
 - e.g., for wordsize of 8, 00000000 and 11111111 are both 0 in 1's complement

Problems with signed mag & 1's C

- Suppose we wanted to have a representation that includes negative numbers where we can apply the **binary addition algorithm** (BAA) to perform the addition:
- 1's complement and signed magnitude won't always work
- E.g., 4-bit word example:

$$\begin{array}{r} +0101 \\ 1101 \\ \hline \end{array}$$

Using binary addition algorithm

Problems with signed mag & 1's C

- Suppose we wanted to have a representation that includes negative numbers where we can apply the **binary addition algorithm** (BAA) to perform the addition:
- 1's complement and signed magnitude won't always work
- E.g., 4-bit word example:

$$\begin{array}{r} 1 \quad 1 \quad 1 \\ + 0101 \\ 1101 \\ \hline 0010 \end{array}$$

Using binary addition algorithm

Problems with signed mag & 1's C

- Suppose we wanted to have a representation that includes negative numbers where we can apply the **binary addition algorithm** (BAA) to perform the addition:
- 1's complement and signed magnitude won't always work
- E.g., 4-bit word example:

		Signed magnitude
1	1	1
0	1	0
1	1	0
1	1	0
<hr/>		
0	0	1

Result should be 0, not 2

Using binary addition algorithm (BAA)

Problems with signed mag & 1's C

- Suppose we wanted to have a representation that includes negative numbers where we can apply the binary addition algorithm to perform the addition:
- 1's complement and signed magnitude won't always work
- e.g., 4 bit example:

	Signed magnitude	1's complement
$\begin{array}{r} 1 \quad 1 \quad 1 \\ + 0101 \\ 1101 \\ \hline 0010 \end{array}$	$\begin{array}{r} + 5 \\ + -5 \end{array}$	$\begin{array}{r} + 5 \\ + -2 \end{array}$
Result should be 0, not 2		Result should be 3, not 2

Using binary addition algorithm (BAA)

Negative Numbers: 2's complement representation

- Non-negative #'s have same form as unsigned (and signed-mag)
- To negate a #, flip all bits and then (using binary add algorithm) add 1
 - Ignore overflow (will only occur when negating 0)
- e.g., wordsize = 4
 - $0010 = 2$ so $1101 + 0001 = 1110 = -2$
 - $0011 = 3$ so $1100 + 0001 = 1101 = -3$
 - $1101 = -3$ so $0010 + 0001 = 0011 = 3$ (can negate negatives too!)
 - $0000 = 0$ so $1111 + 0001 = 0000 = 0$ (0 is unique in 2's complement)
- Note: negation works both ways (going + to - or - to +)
 - The exception: 1 followed by all 0's, e.g., 1000
 - value is -2^{k-1} for wordsize of k, e.g., k=4, value is -8
 - Note: the positive value of 2^{k-1} not expressible with k bits in 2's complement form

Number encodings

	Unsigned	Sign&Mag.	1s comp.	2s comp.
0 0 0	0	+0	+0	+0
0 0 1	1	+1	+1	+1
0 1 0	2	+2	+2	+2
0 1 1	3	+3	+3	+3
1 0 0	4	0	-3	-4
1 0 1	5	-1	-2	-3
1 1 0	6	-2	-1	-2
1 1 1	7	-3	0	-1

8 values

*7 values,
2 zeroes*

*7 values,
2 zeroes*

*8 values,
1 zero*

Number encodings

Signed Mag and 1's complement waste a bit-pattern: 2 reps for 0

	Unsigned	Sign&Mag.	1s comp.	2s comp.
0 0 0	0	+0	+0	+0
0 0 1	1	+1	+1	+1
0 1 0	2	+2	+2	+2
0 1 1	3	+3	+3	+3
1 0 0	4	-0	-3	-4
1 0 1	5	-1	-2	-3
1 1 0	6	-2	-1	-2
1 1 1	7	-3	-0	-1

8 values

*7 values,
2 zeroes*

*7 values,
2 zeroes*

*8 values,
1 zero*

Number encodings

2's complement: smallest negative # has no positive counterpart in representation
The flip-bits-and-add-1 process doesn't successfully negate that number

	Unsigned	Sign&Mag.	1s comp.	2s comp.
0 0 0	0	+0	+0	+0
0 0 1	1	+1	+1	+1
0 1 0	2	+2	+2	+2
0 1 1	3	+3	+3	+3
1 0 0	4	0	-3	-4
1 0 1	5	-1	-2	-3
1 1 0	6	-2	-1	-2
1 1 1	7	-3	0	-1

8 values

*7 values,
2 zeroes*

*7 values,
2 zeroes*

*8 values,
1 zero*

In 2's complement, the binary addition algorithm always works (unless there is overflow)!



- e.g.,

$$\begin{array}{r} \overset{1}{} \overset{1}{} \overset{1}{} \\ + 0101 \quad 5 \\ 1101 \quad -3 \\ \hline 0010 \quad 2 \end{array}$$

Using binary addition algorithm

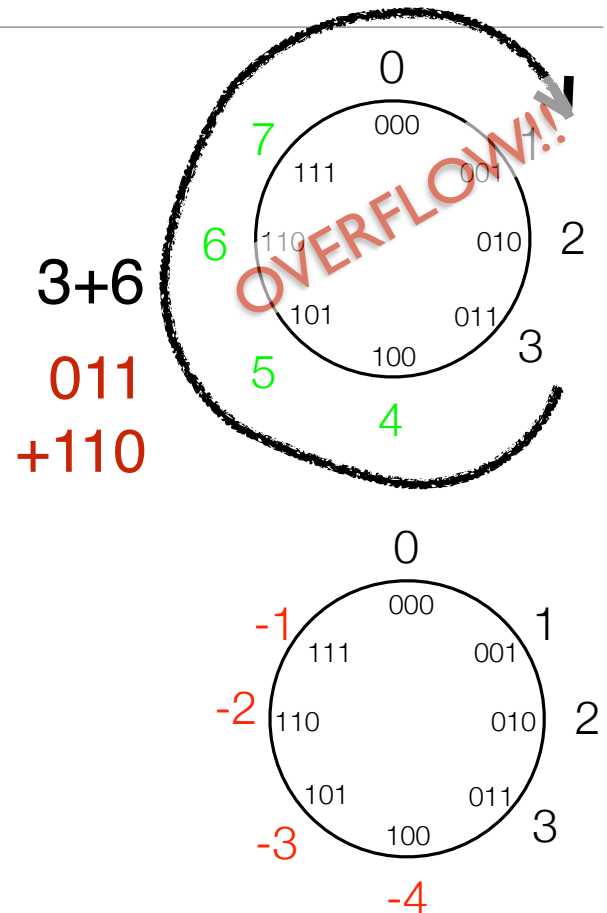
What's special about 2's complement

- Like unsigned, when adding using **BAA**, results are correct mod 2^k (k is word size)
 - Pinwheel perspective: when incrementing, rotate right, when decrementing, rotate left
- In either unsigned or 2's C form, will land on value correct Mod 2^k but actual value not correct if not on pinwheel




What's special about 2's complement

- Think about modular arithmetic (e.g., mod 8)
 - If we choose “unsigned” mod 8 = (0,1,2,3,4,5,6,7)
 - $3 + 6 \bmod 8 = 1$ (i.e., rem (9/8))
 - $3 - 6 \bmod 8 = 5$




What's special about 2's complement

- Think about modular arithmetic (e.g., mod 8)
 - If we choose “unsigned” mod 8 = (0,1,2,3,4,5,6,7)
 - $3 + 6 \bmod 8 = 1$ (i.e., rem (9/8))
 - $3 - 6 \bmod 8 = 5$ (also overflow)

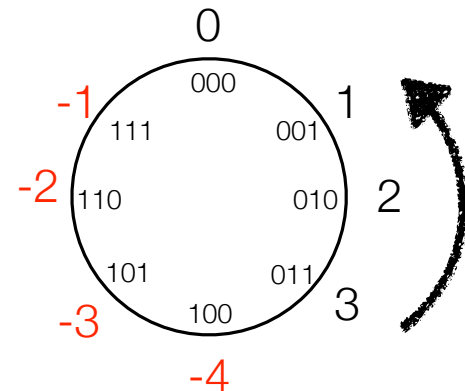


What's special about 2's complement

- But could also choose “2's C” mod 8 = (-4,-3,-2,-1,0,1,2,3)
 - 6 (i.e., 110) would become -2
 - $3 + (-2) \bmod 8 = 3 - 2 \bmod 8 = 1$
 - $3 - (-2) \bmod 8 = 3 + 2 \bmod 8 = -3$


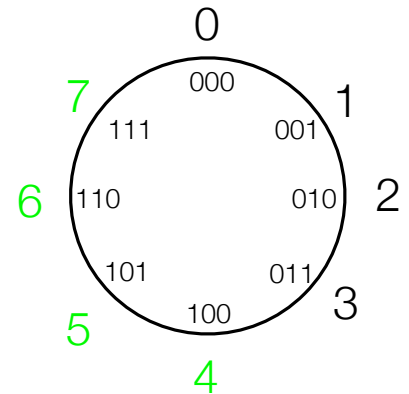


011
+110



What's special about 2's complement

- But could also choose “2's C” mod 8 = (-4,-3,-2,-1,0,1,2,3)
 - 6 (i.e., 110) would become -2
 - $3 + (-2) \bmod 8 = 3 - 2 \bmod 8 = 1$
 - $3 - (-2) \bmod 8 = 3 + 2 \bmod 8 = -3$



Why Flip all bits +1 for 2C works: intuitive “proof”

- Start with any k-bit word X (other than 100...000) and look at it in 2's complement representation
- Let $Y = X$ with all bits flipped
- $X+Y = 11...1111$ (using binary add algorithm)
 - Thus, $X+Y = -1$ in 2's complement rep. (rotated 1 back from 00..0000 on pinwheel)
- Then $X + Y + 1 = 00...0000$ (using binary add algorithm)
- So $Y+1 = -X$ (and recall Y is X with all bits flipped)

Representation v. Operation

k-bit Words & Ranges of various representations

- Given a k-bit word, the range of numbers can be **represented** is:
 - unsigned: 0 to $2^k - 1$ (e.g., k=8, 0 to 255)
 - signed mag: $-2^{k-1} + 1$ to $2^{k-1} - 1$ (e.g., k=8, -127 to 127 [2 vals for 0])
 - 1's complement: same as signed mag (but negative numbers are represented differently)
 - 2's complement: -2^{k-1} to $2^{k-1} - 1$ (e.g., k=8, -128 to 127 [1 val for 0])

Getting representation

- Q: Given an 8-bit wordsize, what is the value of 10001011?

Getting representation

- Q: Given an 8-bit wordsize, what is the value of 10001011?
- A: Is it to be represented in Unsigned, Signed Magnitude, 1's complement or 2's complement?
 - Unsigned: $128 + 8 + 2 + 1 = 139$
 - Signed Mag: $-1 * (8 + 2 + 1) = -11$
 - 1's Complement: the negation of 01110100 = -116
 - 2's Complement: the negation of 01110101 = -117
 - Note: for a given set of bits, when # is negative, 2's complement is 1 less than 1's complement

Representation v. Operation

- We have discussed various **representations** for expressing integers
 - unsigned, signed magnitude, 1's-complement, 2's-complement
- There are also bit-oriented **operations** that go by the same names
 - 1's-complement: flip all bits
 - 2's-complement: flip all bits and (binary) add 1
- Operation can be performed on a number, regardless of representation
 - e.g., let 10111 be a number in signed-magnitude form (value is -7)
 - 2's complement (operation) on 10111 = 01001 (value is 9 in signed-mag form)
- Observe:
 - **2's-complement operation negates a number when number uses 2's-complement representation**
 - **1's-complement operation negates a number when number uses 1's-complement representation**

Automating Subtraction

- Q: Why are we interested in 2's-complement when it seems so less intuitive?
- A: much easier to automate subtraction (i.e., add #'s of opposite sign)
 - e.g., wordsize 6, perform 14 - 21 using signed magnitude representation

$$\begin{array}{r} 001110 \\ - 010101 \\ \hline \end{array}$$

Automating Subtraction

- Q: Why are we interested in 2's-complement when it seems so less intuitive?
- A: much easier to automate subtraction (i.e., add #'s of opposite sign)
 - e.g., wordsize 6, perform 14 - 21 using signed magnitude representation

$$\begin{array}{r} - \quad 001110 \\ \quad 010101 \\ \hline \end{array} \longrightarrow \begin{array}{r} \quad \quad \quad \begin{array}{c} 0 \ 1 \ 10 \\ \text{010101} \end{array} \\ - \quad \quad \quad \begin{array}{c} \text{001110} \end{array} \\ \hline \quad \quad \quad 000111 \end{array}$$

- Signed magnitude has lots of potential “gruntwork”
 - e.g., flip top & bottom, “borrow” from higher order bits, etc.

2's-complement subtraction: make use of BAA

- Just negate subtrahend (bottom # in subtract) and add
- e.g, wordsize 6, perform 14 - 21 using 2's complement representation

$$\begin{array}{r} 001110 \\ - 010101 \\ \hline \end{array}$$

2's-complement subtraction: make use of BAA

- Just negate subtrahend (bottom # in subtract) and add
- e.g, wordsize 6, perform 14 - 21 using 2's complement representation

$$\begin{array}{r} 001110 \quad (14) \\ - 010101 \quad (21) \\ \hline \end{array} \quad \begin{array}{c} \xrightarrow[\text{2's complement operation}]{\text{negate subtrahend}} \\ + \end{array} \quad \begin{array}{r} 001110 \quad (14) \\ + 101011 \quad (-21) \\ \hline \end{array}$$

2's-complement subtraction: make use of BAA

- Just negate subtrahend (bottom # in subtract) and add
- e.g, wordsize 6, perform 14 - 21 using 2's complement representation

$$\begin{array}{rcl} \begin{array}{r} 001110 \quad (14) \\ - 010101 \quad (21) \\ \hline \end{array} & \xrightarrow[\text{2's complement operation}]{\text{negate subtrahend}} & \begin{array}{r} 111 \\ 001110 \quad (14) \\ + 101011 \quad (-21) \\ \hline 111001 \quad (-7) \end{array} \end{array}$$

Apply binary addition algorithm

$$X = 111001, -X = 000111 = 7, X = -7$$

Detecting Overflow

- Q: How to tell if the result of a computation overflows (i.e., result cannot be expressed within the wordsize constraint)

- e.g., 4-bit words, unsigned: $1110 + 1010$ ($14 + 10$)

- result is 24, cannot be expressed in 4-bit unsigned (only values 0-15)


- Hence, overflows

- Detection in unsigned easy:

- Addition: if final carry = 1, then overflow, else not

- Subtraction: subtrahend is bigger # (result negative), then overflow

Indication of overflow
for unsigned


$$\begin{array}{r} 111 \\ 1110 \\ 1010 \\ \hline 1000 \end{array}$$

Overflow detection in 2's-complement is easy

- If final two carries match, then no overflow
- If differ, then overflow
- e.g., wordsize = 4

$$\begin{array}{r} 5+1=6 \\ \textcolor{red}{00} \\ + 0101 \\ \hline 0001 \\ + 0110 \\ \hline \end{array}$$

$$\begin{array}{r} -5 + -3 = -8 \\ \textcolor{red}{11} \\ + 1011 \\ \hline 1101 \\ + 1101 \\ \hline 1000 \end{array}$$

$$\begin{array}{r} -2 + 7 = 5 \\ \textcolor{red}{11} \\ + 1110 \\ \hline 0111 \\ + 0111 \\ \hline 0101 \end{array}$$

$$\begin{array}{r} -2 + -7 = -9 \\ \textcolor{red}{10} \leftarrow \text{overflow} \\ + 1110 \\ \hline 1001 \\ + 0111 \\ \hline \end{array}$$

$$\begin{array}{r} 7+7=14 \\ \textcolor{red}{01} \leftarrow \text{overflow} \\ + 0111 \\ \hline 0111 \\ + 0111 \\ \hline 1110 \end{array}$$

Remainder of Course: We use unsigned and 2's complement: It's what computers usually use for their representation

Where have you used 2's complement (w/o knowing it?)

```
int main() {  
    int x = 47;  
    int y = -50;  
    unsigned int z = 50;  
    printf("%d %d %u\n", x, y, z);  
}
```

Where have you used 2's complement (w/o knowing it?)

```
int main() {  
    int x = 47;  
    int y = -50;  
    unsigned int z = 50;  
    printf("%d %d %u\n", x, y, z);  
}
```

- The computer represents
 - x and y as signed 2's complement
 - z as an unsigned

Final thoughts on Overflow

- Can't stress enough that **overflow means the result cannot be represented within the word size** (not necessarily that the result is too big)!
- Example: suppose I choose a (weird) way to map 2-bit words to values as follows:

Two-bit word	Associated Value
00	1
01	3
10	6
11	7

Final thoughts on Overflow

- Can't stress enough that **overflow means the result cannot be represented within the word size** (not necessarily that the result is too big)!
- Example: suppose I choose a (weird) way to map 2-bit words to values as follows:

Two-bit word	Associated Value
00	1
01	3
10	6
11	2

$$00 + 00 = 11 \quad (\text{i.e., } 1 + 1 = 2)$$

$$00 + 01 = ??$$

1 + 3 should equal 4, but no 2-bit combo represents 4

Final thoughts on Overflow

- Can't stress enough that **overflow means the result cannot be represented within the word size** (not necessarily that the result is too big)!
- Example: suppose I choose a (weird) way to map 2-bit words to values as follows:

Two-bit word	Associated Value
00	1
01	3
10	6
11	2

$00 + 00 = 11$ (i.e., $1 + 1 = 2$)

$00 + 01 = ??$

$1 + 3$ should equal 4, but no 2-bit combo represents 4

Thus, for this representation, $00+01$ causes overflow

Floating Point

Need a bigger range? Floating Point Representation

- Change the encoding.
- Floating point (used to represent very large numbers in a compact way)

- A lot like scientific notation: $-7.776 \times 10^3 = -7776 = -6^5$

mantissa
(a.k.a. fraction)

Note: mantissa always in form X.XX...
(one digit before the '.')

- But for this course, think binary:

$$-1.10 \times 2^{0111} (= -1.5 * 2^7)$$


$$1.01 \times 2^{-0111} (= 1.25 * 2^{-7})$$

Note: in proper form, for binary, mantissa always 1.XX...
(one digit before the '.' and digit always a 1)

The only exception: $0 = 0.0 \times 2^0$


Standard Forms for Floating Point Numbers

- How to represent a floating point # within the confines of a 32-bit word
- The bits of the word are separated into different **fields**




- IEEE 754 standard specifies
 - which bits represent which fields (bit 31 is sign, bits 30-23 are 8-bit exponent, bits 22-00 are 23-bit fraction)
 - how to interpret each field

IEEE 754 Floating Point description



- **Sign**: 0 = positive #, 1 = negative # (like signed magnitude)
- **Exponent**: **unsigned** with a **bias** of 127
 - exponent value = unsigned binary rep of 8 bits - 127
- **Fraction**: recall fraction always of form 1.XXXXXXX
 - leave off the '1', just represent the "XXXXXXX" part



$$= (-1) \times 1.101_{(2)} \times 2^{13-127} = -1.625_{(10)} \times 2^{-114} = -7.82409 \times 10^{-35}$$

Approximation of 0

- Because representation always of form $\pm 1.XXXX \times 2^{yyyy}$, cannot represent true 0
- Note: All bits set to 0 equals 1.0×2^{-127} , a very very small number, effectively 0

Bias and Comparing floats

- Bias allows exponents between -127 (very small) and 128 (very large)
- Q: Why bias instead of 2's-complement or unsigned magnitude?
- A: Easy comparison between two floats, A & B, for which is larger
 - **Step 1:** Check signs. A+ and B-, return $A > B$. A- and B+, return $A < B$
 - **Step 2** (A and B same sign): Check exponents
 - (A+ and A.exp > B.exp) or (A- and A.exp < B.exp), return $A > B$
 - (A- and A.exp > B.exp) or (A+ and A.exp < B.exp), return $A < B$
 - **Step 3** (A and B same sign, same exponents): check fractions
 - (A+ and A.frac > B.frac) or (A- and A.frac < B.frac), return $A > B$
 - (A- and A.frac > B.frac) or (A+ and A.frac < B.frac), return $A < B$
 - **Step 4** (A and B same sign, same exponents, same fraction): return $A = B$
- Observation: when bits are ordered sign, exponent, magnitude, process same as comparing 2 signed-magnitude numbers

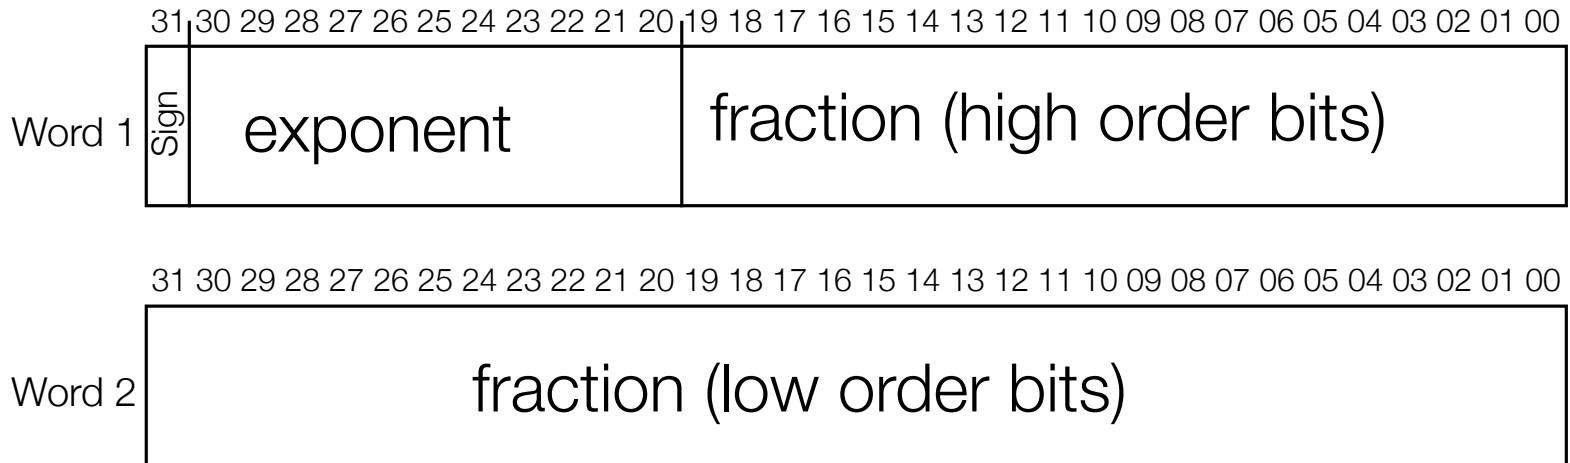
Which IEEE 754 FP # is biggest?



- All + #'s, $B > A$ (larger exponent), $B > C$ (same exponent, larger fraction)
- What if #'s were simply viewed as signed-mag?
 - Same result: $B > C > A$

IEEE 754 64-bit (double) precision

- Described within 32-bit architecture as **two 32-bit words**



- 11-bit exponent with bias of 1023
- fraction is 52 bits long (20 higher order bits in Word 1, remaining 32 lower order bits in Word 2)

Underflow

- When the **magnitude** of the value is **too small** to be described by the representation
- e.g., in IEEE 754 floating point:
 - $X = 1 \times 2^{-100}$ can be represented by the standard (what is the set of bits in the exponent field?)
 - However, X^2 , which equals 2^{-200} is too small to be represented
 - smallest possible exponent is -127
 - Hence, attempts to compute $X \times X$ results in an underflow

End of Lecture

- Important take-aways for rest of course:
 - High-level view of computer (digital perspective): notion of word size
 - Binary #'s: negative representation (2's complement)
 - adding, subtracting, overflow, overflow detection